

Introduction to Artificial Intelligence

LECTURE 3: Faster Complete Search

Overview

- Dynamic Programming
- Minimum Spanning Trees
- Prim's Algorithm
- Shortest Path Problem
- Dijkstra's Algorithm
- A* Search

Dynamic Programming

- Synthesises complete solution from partial solutions

Synthesisability Condition: (when) complete solutions *can* be built up from partial solutions.

- Bottom-up:
 - starts from a null solution
 - solves a base case (or cases)
 - composes larger partial solutions, until done
- Often allows for a trade-off: use extra storage for intermediate results; avoid re-computing intermediate results

Fibonacci Sequence

Fibonacci problem clearly is synthesisable:


$$f(n) = f(n - 1) + f(n - 2)$$

So, to compute FIB more quickly, we can save intermediate partial solutions rather than re-compute them:

- (1) INT FIB (INT n)
- (2) INT $a[0 : n]$
- (3) $a[0] \leftarrow 0; a[1] \leftarrow 1;$
- (4) IF $n \leq 1$ RETURN n
- (5) ELSE FOR $i \leftarrow 2$ TO n
- (6) $a[i] \leftarrow a[i - 1] + a[i - 2];$
- (7) RETURN $a[n];$

TSP - Dynamic Programming

Consider a simple (symmetric) TSP problem:



where $L(i, \text{neighbour}(i)) = (\text{symmetric})$ edge weight between the two nodes

TSP - Dynamic Programming

Let $g(i, S) =$


the length of the shortest path from i to 1
visiting each city in S exactly once (where S is
a sub-graph of some graph) and possibly also
going through one or more of any cities in the
graph which are not in S

Goal: Find $g(1, V - \{1\})$ for some graph, V

Synthesizability Condition:

$$(*) \quad g(i, S) = \min_{j \in S} \{L(i, j) + g(j, S - \{j\})\}$$

TSP - Dynamic Programming



The circuits and their costs are

$\langle 1234 \rangle$	=	18
$\langle 1243 \rangle$	=	19
$\langle 1324 \rangle$	=	23
$\langle 1342 \rangle$	=	19
$\langle 1423 \rangle$	=	23
$\langle 1432 \rangle$	=	18

$$\begin{aligned}
 g(1, V - \{1\}) &= \min_{j \in V - \{1\}} \{L(1, j) + g(j, V - \{1, j\})\} \\
 &= L(1, 2) + g(2, V - \{1, 2\}) \\
 &= L(1, 2) + g(2, \{3, 4\}) \\
 &= 2 + \min(16, 17) \\
 &= 18
 \end{aligned}$$

TSP - Dynamic Programming

Idea: solve the base case; figure out how to synthesise sub-answers into larger answers.

Base case: $S = \emptyset$

$$g(1, \emptyset) = 0$$

$$g(2, \emptyset) = 2$$

$$g(3, \emptyset) = 6$$

$$g(4, \emptyset) = 7$$

$|S| = 1$

$$g(2, \{3\}) = 10$$

$$g(2, \{4\}) = 13$$

$$g(3, \{2\}) = 6$$

$$g(3, \{4\}) = 12$$

$$g(4, \{2\}) = 8$$

$$g(4, \{3\}) = 11$$

And iterate ...

TSP - Dynamic Programming Algorithm

```
(1) INT TSP ( $L[*,*], P[*,*]$ )
(2) INT  $g[*,*], n \leftarrow \dim(L)$ ;
(3) FOR  $i = 1$  TO  $n$   $g[i, \emptyset] \leftarrow L[i, 1]$ ;
(4) FOR  $k = 1$  TO  $n - 2$ 
(5)   [FOR ALL  $S \subseteq V - \{1\}$  S.T.  $|S| = k$ 
(6)     [FOR  $i \notin S \cup \{1\}$ 
(7)       [ $g[i, S] \leftarrow \min_{j \in S} (L[i, j] + g[j, S - \{j\}])$ ;
(8)        $P[i, S] \leftarrow \text{THAT } j$ ];]]
(9)  $g[1, V - \{1\}] \leftarrow \min_{j \in V - \{1\}} (L[1, j] + g[j, V - \{1, j\}])$ ;
(10)  $P[1, V - \{1\}] \leftarrow \text{THAT } j$ ;
(11) RETURN  $g[1, V - \{1\}]$ ;
```


Return value is the min length; returned P reports the min path.

TSP : DP VS Complete Search

Cities	DP [2^{n-3}]	Brute Force [$(n-1)!/2$]
1	1	1
2	2	3
3	4	12
4	8	60
5	16	360
6	32	2520
7	64	20160
8	128	181440
9	256	1814400
10	512	19958400
11	1024	239500800
12	2048	3113510400
13	4096	43589145600
14	8192	653837184000
15	16384	10461394944000
16	32768	177843714048000
17	65536	3201186852864000
18	131072	60822550204416000

Spanning Trees


- Definition 1: Connected Graph
 - A **connected** graph is one in which there is a path from each node to every other node.



- Definition 2: Tree
 - A **tree** is an acyclic, connected graph
- Definition 3: Spanning Tree
 - A **spanning tree** is a tree that contains all the nodes of the original graph.

Minimum Spanning Tree

- The minimum spanning tree problem is to find a spanning tree which minimizes the sum of weights on a spanning tree.



Minimum Spanning Tree Algorithm

(1) $V' := \text{random}(V)$; [CHOOSE THE ROOT]

(2) $E' := \emptyset$

(3) WHILE $V' \neq V$ DO

(4) $[E' := E' \cup E(\text{nearest}(V', V \setminus V'))];$

(5) $V' := V' \cup \text{nearest}(V', V \setminus V')]$

- *nearest* return the nearest node not in V'
- $E(\text{nearest})$ returns the edge to that node

Prim's Algorithm


- Outer loop executes $n-1$ times
- Let us examine *nearest*

```
(1)  $min := \infty; next := null;$ 
(2) FOR  $v \in V - V'$ 
(3)   FOR  $w \in V'$ 
(4)     [IF  $d(v, w) < min$ 
(5)       THEN [ $next := v; min := d(v, w)$ ]];
(6) IF  $next \neq null$  RETURN  $next$  ELSE FAIL;
```

- Inner loop (e.g., (4) and (5) above) is executed $(n - k)k$ times, where $k = |V'|$.
- This is maximized when $k = n/2$ (differentiate $(n - k)k$ wrt k and set to zero), when it executes $n^2/4$ times.
- So, this algorithm is $O(n^3)$

Shortest Path Problem

- Problem formulation
 - For any two nodes $v, w \in V$ find the shortest path between them.
- Dijkstra's Algorithm (vaguely analogous to Prim's)
 - Expand the partial path with the lowest cost until the complete solution has been built.
- Example: What is the shortest path from node 1 to 6?




Dijkstra's Algorithm

```
(1) open  $\leftarrow$  initial; closed  $\leftarrow$   $\langle \rangle$ ;  
(2) LOOP: IF open =  $\langle \rangle$  FAIL;  
(3) v  $\leftarrow$  pop(open);  
(4) push(v, closed);  
(5) IF goal(v) THEN RETURN path(initial, v)  
(6) ELSE open  $\leftarrow$  sort(children(v), open);  
(7) GO LOOP;
```

- *initial* contains the starting node
- *goal* is a test for the target node
- The sort operation sorts nodes by retrospective pathcost – lowest cost of the partial path to the node being added to *open*.

Dijkstra's Algorithm

Example:



At step 2 the open and closed lists will be:

$\langle 1 \rangle$	$\langle \rangle$
$\langle 324 \rangle$	$\langle 1 \rangle$
$\langle 245 \rangle$	$\langle 13 \rangle$
$\langle 45 \rangle$	$\langle 132 \rangle$
$\langle 56 \rangle$	$\langle 1324 \rangle$
$\langle 6 \rangle$	$\langle 13245 \rangle$
$\langle \rangle$	$\langle 132456 \rangle$

Dijkstra's Algorithm


- Dynamic programming synthesizing approach (complete search)
 - Each partial path is guaranteed to have minimum cost when added; therefore, when the final path is added it will have minimum cost.
 - $O(n^2)$
- Greediness?

A* Search Algorithm

- Generalization of Dijkstra's algorithm by considering heuristic evaluation
 - Sorting based on
 - $c(x)$ = retrospective cost from i to x
 - $h(x)$ = prospective heuristic cost from x to the nearest goal
 - $eval(x) = c(x) + h(x) = A^*$ eval of the path
 - $best(X) = \min_{x \in X} \{eval(x)\}$
- Note: this is “heuristic” only in the sense of using a heuristic evaluation function. Not in the original sense of a hit-or-miss heuristic: it is *complete*!

A* Search Algorithm

E.g., suppose our heuristic h is perfect. Then



At step 2 the open and closed lists will be:

$\langle 1 \rangle$	$\langle \rangle$
$\langle 324 \rangle$	$\langle 1 \rangle$
$\langle 524 \rangle$	$\langle 13 \rangle$
$\langle 624 \rangle$	$\langle 135 \rangle$
$\langle 24 \rangle$	$\langle 1356 \rangle$

Qualitatively (check visually on larger graphs):

- Dijkstra's algorithm spreads search out in all directions
- A^* (with a good h) focuses search along a promising set of paths (in a "finger")

A* Search Algorithm

- Quality of the algorithm depends upon the heuristic evaluation function h .
- The main issue is to generate an *admissible* heuristic.
 - h is **admissible** iff $\forall S \ h(S) \leq \text{dist}(S, G)$
- Theorem: A* is complete iff. its heuristic is admissible
- Given admissibility, efficiency of the algorithm improves as $h(x)$ converges on the true cost to goal.

Both Dijkstra's and A* Algorithms are not Greedy

- Informal 'Proof'
 - Greedy algorithms make the best looking single step from wherever they are now.
 - In A* the best looking single step is given by $h(x)$; however, A* takes that only if $h(x) + c(x)$ happens to be best.
 - Dijkstra's takes the best looking next step only if the cost to it plus total retrospective cost happens to be best.
 - Incidentally, Dijkstra could only be greedy if A* is greedy, since it *is* A* with $h(x) = 0$ everywhere.
- A* is systematic and complete (but, given a good h , not exhaustive).

3 Kinds of Search

1. *Complete (Optimal) Search*: Guaranteed to find the goal (optimum), if there is one. E.g.,
 - Exhaustive search
2. *Heuristic Search*: Not necessarily guaranteed to find the (optimal) answer. E.g.,
 - Greedy search
3. *Stochastic Search*: probabilistic search through the state space. Not guaranteed, hence a (sub)kind of heuristic search. E.g.,
 - Genetic algorithm, simulated annealing